Knowledge & the web: exercise session 6

Katrien Laenen   Gust Verbruggen   Ward Schodts

November 25, 2015

In this report we discuss which algorithms we tried to find a good classification model that predicts if a passenger of the Titanic will survive or not.

We use the data from `titanic_train.csv` to train our model and perform cross-validation for performance assessment. The models used are a **rule set** and a **decision tree**. The rest of this document is structured as follows. First we describe the preprocessing steps used in order to prepare the data for better classification. Next, we describe the learning process for both models. Finally, we compare results for both learning techniques, based on the cross-validation results.

1 Preprocessing

In order to achieve better classification, we preprocess the data. This is done in two phases: feature selection to only preserve relevant features and feature engineering in the form of a discretisation step.

Feature selection   Features that we do not think are relevant are `Name`, `Ticket`, `PassengerId`, `Embarked` and `Cabin`. They are removed.

Discretisation   The two numerical values are discretised in accordance to some criterium. The `Age` feature into five age classes: child (< 9 years), teen (10 to 15 years), young (15 to 24 years), adult (25 to 59 years) and elderly (> 59 years). To discretise `Fare`, we take a look at a histogram plot, as seen in Figure 1. There are a lot more low fares. The three leftmost bars get their own bin. The chosen discretisation is: 0-10, 10-20, 20-30, 30-100 and 100+.

2 Rule learning

One thing we tried was to learn classification rules from the training data. Rule learning algorithms come with different options for several aspects of the algorithms. In the exercise session we experimented with two of these.

Figure 1: Histogram for the `Fare` feature

First, for choosing the conditions of a classification rule, there are two options: one is to add to the rule the condition which increases the accuracy the most, another is to add to the rule the condition which increases the information gain the most. The accuracy criterion option resulted in the following rule set:

- if Parch > 5.500 then 0 (1 / 0)

- if Fare $\geq$ 10 and Sex = male then 0 (170 / 25)

- if Sex = female then 1 (63 / 197)

- else 0 (171 / 61)

The information gain criterion option resulted in this rule set:

- if Sex = male and Fare $\geq$ 10 then 0 (170 / 25)

- if Sex = female then 1 (64 / 197)

- if Fare = 10-20 and Parch $\leq$ 0.500 then 0 (66 / 7)

- else 0 (111 / 54)

A second aspect of rule learning is the desired pureness, which is the minimum ratio of the major class in a covered subset in order to consider the subset pure. The standard value for this is 0.9. We experimented with this value set to 0.8 and 1. For the desired pureness set to 0.8 and the accuracy criterion we got the following rule set:

2

- if Parch > 5.500 then 0 (1 / 0)

- if Fare ≥ 10 then 0 (189 / 47)

- if Sex = female then 1 (44 / 175)

- else 0 (171 / 61)

This rule set is very similar to the rule set obtained with the accuracy criterion option and desired pureness of 0.9. The only difference is that in the second rule the *Sex = male* condition now disappeared from the condition part. For the desired pureness set to 0.8 and the information gain criterion we got the following rule set:

- if Sex = male then 0 (360 / 93)

- else 1 (57 / 177)

When comparing this rule set with the rule set obtained for the information gain criterion with the desired pureness set to 0.9, we see that this rule set is a great simplification of the other. For the desired pureness set to 1 and the accuracy criterion we got exactly the same rule set as when setting the desired pureness to 0.9 for that criterion. The same happened for the desired pureness set to 1 and the information gain criterion, which also resulted in the same rule set as when setting the pureness to 0.9 and keeping the information criterion.

Next to the criterion and desired pureness, there are three other parameters for rule learning: sample ratio (which specifies the sample ratio of training data used for growing and pruning), minimal prune benefit (which specifies the minimum amount of benefit which must be exceeded over unpruned benefit in order to be pruned) and use local random seed (which indicates if a local random seed should be used for randomization). We did not experiment with these parameters, but kept the standard values for these, which were 0.9, 0.25 and no respectively.

From these results we can clearly see that Sex, Fare and Parch are features which are important for this task. That Sex and Fare were important features was to be expected. When the Titanic collided with an iceberg, women and children and people of the higher ranks were evacuated first. This is reflected in the data as we see that women and people which payed a higher fare are more likely to have survived. That the number of parents and children is also an important feature can be explained as follows. People in need want to stay close to their loved ones at all costs. But how larger a group, how more difficult it was to get on a lifeboat together. Also, little children and old people need much more assistence to even reach a lifeboat. As such, how more children and parents someone had aboard, how less likely that person was to survive.

3 Decision trees

Apart from rule set learning, we also experimented with decision trees. Since we are working with a mix of nominal and numeric data, this is a suitable model, because each node can use its own *rule* for splitting on an attribute. We will describe how we tweaked the learning parameters in order to achieve better results. The initial tree is shown in Figure 2.
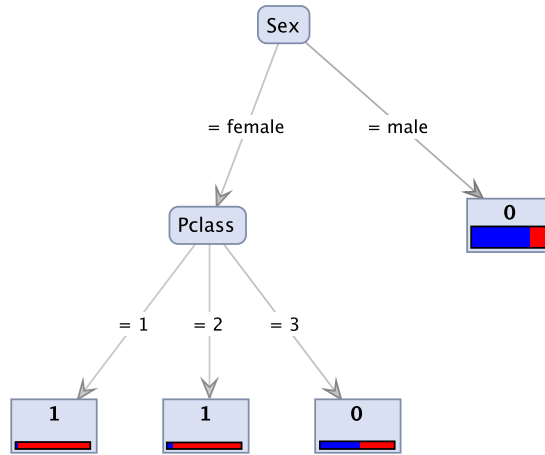

Figure 2: Tree learned with default parameters

The most obvious observation about this initial model is that female passengers with third class tickets are not yet classified very accurately, yet the tree did not split on it. In order to improve on this, we can adjust the splitting criterion to *accuracy*. It will split nodes such that the accuracy of the tree increases. This results in a very sparse and large tree, which is not shown for exactly that reason. By changing the maximal depth of the tree, we can prevent it from getting too large. A maximal depth of 4 was found to provide good results. The result is shown in Figure 3. As we can see, this achieved the desired result of splitting the undecided node in four new nodes, three of which perform clearly better. The final step is improving prediction for males. While its accuracy is quite high, the recall for surviving males is zero. This probably happens because the gain in criterion is not sufficiently high for the pre-pruning step to split the node. By lowering the minimum gain during pre-pruning, we enable the algorithm to split easier on nodes. Because the `Fare` node is already split until max depth, we expect that only the rightmost node (males) will be further split by this setting. Figure 4 confirms this assumption and shows the final model.

This model is pretty intuitive. For females, Class is the most important deciding factor. Amongst the lowest class women, the ones that payed most for their ticket surprisingly have the lowest survival rate. For males, Age is

Figure 3: Tree learned with *accuracy* splitting criterion and maximum depth of 4.
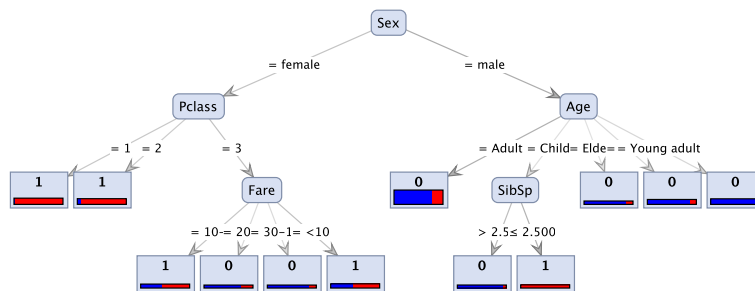

Figure 4: Tree learned with *accuracy* splitting criterion, maximum depth of 4 and minimum information gain of 0.01.

the most important deciding factor. Children have a far higher survival rate than men.

4 Rule learning vs. Decision trees

We conclude this report with a brief summary of our results and our conclusion about the methods. To decide if either Rule learning is better than Decision trees we look at three performance measures, namely: **accuracy**, **precision** and **recall**.

With *accuracy* we mean the proportion of true results (both true positives and true negatives) among the total number of cases examined. For the Titanic case this would mean:

$$accuracy = \frac{\#\text{true dead} + \#\text{true alive}}{\#\text{true dead} + \#\text{false dead} + \#\text{true alive} + \#\text{false alive}}$$

Precision is the fraction of retrieved instances that are relevant:

$$precision = \frac{\#\text{true alive}}{\#\text{true alive} + \#\text{false alive}}$$

Finally *recall* (or sensitivity) is the fraction of relevant instances that are retrieved:

$$recall = \frac{\#\text{true alive}}{\#\text{true alive} + \#\text{false dead}}$$

In following table we give the accuracy, precision and recall of our best configuration from both rule learning and decision trees. With best we mean the highest **accuracy**, **precision** and **recall**.

	Accuracy	Precision	Recall
Rule learning	78.99%	80.59%	65.87%
Decision trees	81.46%	81.37%	71.26%

Figure 5: Our results

From these scores it is obvious to see that the clustering with a decision tree is a lot better. With our configurations decision trees outperforms rule learning in every performance measure.

We can conclude that a **decision tree** works better in this specific case.

5 Further experiments with evaluation methods

In our previous evaluations we used the regular cross-validation. For session 7 we experimented with *leave_one_out* cross-validation (this is just a setting in RapidMiner). Actually it's going to leave out one example of the training set. This example is used to test the model that is build, using all the other examples. This is done for every element in the dataset. This way this is very computational expensive.

Apart from leave_one_out we experimented with *split validation*. This method splits the data at random in a training part and a test part. One can specify the ratio test/training data in RapidMiner. We used the standard value here: 0.7.

In the following tables we give a short summary of our results on the best model from the previous section.

	Accuracy	Precision	Recall
Cross-validation	78.99%	80.59%	65.87%
leave_one_out	78.01%	75.09%	68.62%
Split validation	76.64%	75.00%	62.79%

Figure 6: Results for Rule Learning

	Accuracy	Precision	Recall
Cross-validation	81.46%	81.37%	71.26%
leave_one_out	82.47%	82.35%	72.66%
Split validation	82.46%	82.35%	72.66%

Figure 7: Results for Decision Trees

We can conlude that the new evaluation methods give worse results for the rule learning approach. In contrary the results for decision trees became even better.